

re

# Founding Applied AI Engineer

**Kay.ai**

ID: SRN2025-10561

**Founding Applied AI Engineer**

Kay.ai • Full Time • New York,  
NY • Hybrid • \$140k - \$180k +  
Equity

**About the Company:**

Kay.ai is on a mission to eliminate uninspiring work, starting with repetitive data entry in insurance. The company is backed by top investors in B2B, RPA, and Insurance, with a stacked team of engineers and researchers from Microsoft, Amazon, Twitter, and Relyance. Kay.ai is pioneering the use of LLM-first AI applications to automate complex workflows and improve efficiency in insurance and beyond.

**Roles and Responsibilities:**

- Build LLM-first AI products and features that have real-world usage.
- Develop browser automation AI agents for data extraction, classification, and decision-making.
- Design robust prompting systems, fine-tune models, and create evaluation frameworks for AI systems.
- Work with a modern AI tech stack, including Python, JavaScript, TypeScript, OpenAI, Hugging Face, dspy, ollama, and Docker.
- Deploy serverless, scalable AI-powered systems, leveraging microservices and cloud-based architecture.
- Work closely with the founding team to shape the company's AI strategy and roadmap.
- Handle unstructured document processing (OCR, PDFs) and develop solutions for structured and unstructured data workflows.

**Job Requirements:**

- 3-6 years of software engineering experience with a strong focus on LLM and AI applications.
- Recent hands-on experience (within the last 1-2 years) building LLM-first products that have received real usage.
- Experience working at high-growth AI startups (Seed - Series B).
- Expert proficiency in Python, including dspy, Hugging Face, OpenAI/Gemini APIs, Ollama.
- Intermediate experience in JavaScript and TypeScript for full-stack AI application development.
- Strong foundational software engineering knowledge, including serverless architecture, Docker, TDD, and microservices.
- Experience working with AI agents (e.g., CrewAI) or unstructured document processing (OCR, PDFs).
- Stays up-to-date with the latest AI models, benchmarks, and open-source contributions.
- Passion for 0-to-1 development, building new products from the ground up.
- Computer Science degree from a top-tier university, such as MIT, Stanford, CMU, UC Berkeley, Harvard, UIUC, University of Toronto, etc.

**✗ Do NOT Apply If You:**

- Require visa sponsorship (H1B, TN, etc.).
- Have a history of frequent short tenures (job hopping).
- Have only worked in big corporations (Uber, Intel, etc.) without startup experience.
- Come from an IT consulting background (Infosys, Tata, Capgemini, Cognizant, Wipro, etc.).
- Graduated from coding bootcamps (Full Stack Academy, Hack Reactor, etc.).
- Have fake or misrepresented profiles.

**Interview Process**

1. Intro Call (30 min) – Initial discussion on qualifications and non-negotiables.
2. Technical Round (1 hour) – Hands-on technical evaluation with an engineer.
3. Pair Programming Interview (30 min) – Implement a common task using a new framework.
4. On-Site in NYC (1-2 days, Paid) – Final round assessment, willing to fly out candidates.
5. Offer Extended